

EDUCATION

- Indian Institute of Technology, Bombay

MS by Research, Computer Science and Engineering GPA: 9.76/10 (Department Rank 2) Aug'21 – June'24
Advisors: Prof. Pushpak Bhattacharyya & Prof. Harshad Khadilkar

Recipient of the “**Winifred B. Fernandes Research Excellence Award**” for my MS Thesis.

PUBLICATIONS & PREPRINTS

(* = EQUAL CONTRIBUTION)

- One Prompt To Rule Them All: LLMs for Opinion Summary Evaluation
Swaroop Nath*, Tejpalsingh Siledar*, Pushpak Bhattacharyya, *et al.* ACL'24
□ paper  
 - Replaces Supervision: Query focused Summarization using Deep Reinforcement Learning
Swaroop Nath, Pushpak Bhattacharyya, and Harshad Khadilkar EMNLP'23
□ paper  
 - Transformers are Expressive, But Are They Expressive Enough for Regression?
Swaroop Nath, Harshad Khadilkar, and Pushpak Bhattacharyya arxiv  
 - Leveraging Domain Knowledge for Efficient Reward Modelling in RLHF: A Case-Study in E-Commerce Opinion Summarization.
Swaroop Nath*, Tejpalsingh Siledar, Pushpak Bhattacharyya, *et al.* arxiv  

RESEARCH EXPERIENCE

Pre-Doctoral Researcher | Google DeepMind

July'24 – Present

Hosts: Dr. Praneeth Netrapalli, Dr. Nitish Gupta, & Dr. Prateek Jain

- Implemented a novel training-free, clustering-based *Efficient Attention algorithm* that led to $\sim 40\%$ reduction in latency while maintaining text generation quality.
 - Developed a novel *Efficient Reasoning algorithm* which reduces verbosity in frontier reasoning models by $\sim 50\%$ while maintaining reasoning quality.

AI/ML Research Intern | LinkedIn

May'23 – July'23

Host: Smit Marvaniya

- Developed a *novel* algorithm for Spam detection in LinkedIn comments, identifying the relevance of context (neighboring comments and post), which led to $\sim 20\%$ quality improvement of the overall pipeline.

SCHOLASTIC ACHIEVEMENTS & NEWS

RESEARCH PROJECTS

Efficient Reward Modelling for RLHF by Leveraging Domain Knowledge

Oct'23 – June'24

Advisors: Prof. Pushpak Bhattacharyya & Prof. Harshad Khadilkar

- Developed a novel Reward Modelling approach for Reinforcement Learning from Human Feedback, which reduced preference data requirement by **20 \times** , while still achieving alignment with humans when used in RLHF.
 - Released two datasets for Opinion Summarization: **PROMPTOPINSUMM** – SFT data, & **OPINPREF** – Preference data.

June'23 – Sept'24

Advisors: Prof. Pushpak Bhattacharyya & Prof. Harshad Khadilkar

- Experimented with Hierarchical Text Generation for QfS, motivated by the underlying textual structure present within summaries.

Reinforcement Learning based Query focused Summarization (QfS) May'22 – June'23

Advisors: Prof. Pushpak Bhattacharyya & Prof. Harshad Khadilkar

- Developed a novel Reinforcement Learning algorithm for QfS – led to **37.2%** improvement (automatic metric) and **19%** improvement (human evaluation).
- Contributed a *gold-standard* dataset (**250 samples**) which tackles Topic Centralization, for analysis of QfS models.

Passage Embedding using Siamese BERT Oct'22 – June'23

Advisors: Prof. Pushpak Bhattacharyya & Prof. Harshad Khadilkar

- Developed a novel *Passage Embedding mechanism* inspired by *Cluster Hypothesis*, which was used in a reward function for RL based QfS training – led to **2.21** point improvement (ROUGE) over the next best reward function.
- Contributed a *new* dataset (~**8 million instances**) for pre-training the Passage Embedding model.

Open Domain Aspect-Based Sentiment Analysis (ABSA) Aug'21 – May'22

Advisor: Prof. Pushpak Bhattacharyya

- Created a new dataset for ABSA from Yelp reviews, motivated by the lack of diversity of domains in existing datasets. It covered **111 domains**, which led to ~ **18%** improvement for models in *open-domain* performance.

SELECTED COURSE PROJECTS

Beyond BranchFormer: What re-stitching the network does! Jan'23 – May'23

Instructor: Prof. Preethi Jyothi

- Patched the *faulty* dropout implementation in the original BranchFormer implementation, and experimented with both *Attention dropout* and *CG-MLP dropout*.
- Tested Linear Attention (from BigBird), achieving a CER of **4.97** (vs **4.22** in the original FastFormer Attention).

Caption Consistency Detection Jan'22 – May'22

Instructor: Prof. Pushpak Bhattacharyya

- Developed a CNN+Transformer-based model for the *novel* task: **Caption-Image Consistency Detection**.
- Implemented a custom **near-miss-near-hit** based sampling technique to mine negative samples for training, which helped the model achieve an F1-score of **89%**.

DOODLERGAN: Exploring GANs for Doodling July'21 – Nov'21

Instructor: Prof. Preethi Jyothi

- Reimplemented DoodlerGAN, a StyleGAN2 inspired recurrent part generator, for the Creative Birds dataset, which led to **75%** more diverse generations than the vanilla StyleGAN2 model.
- Explored the impact of ProGAN style progressive generation on the *convergence time* and *sample efficiency*.

SERVICE

Reviewer for A/A* conferences Feb'24 – Present

- Served as a reviewer for several A/A* CL conferences (ACL, EMNLP, EACL, etc.), and a co-reviewer in NeurIPS'24.
- Nominated as a **great reviewer** for 3 reviewed papers (out of 4) for ACL'25.

Web Chair | CODS-COMAD Conference Sept'22 – Jan'23

- Coordinated with the organizing team to keep the conference website up-to-date for all information to be disseminated to the authors; and smooth running of the conference proceedings.

System Administrator | DGX Servers @ IoE-DIS, IIT Bombay Sept'22 – Dec'23

- Served as the **system administrator**, enabling smooth operation of the DGx-A100 servers at IIT Bombay.

Public Relations Manager | CFILT Lab, IIT Bombay Feb'22 – June'24

- Tasked with handling the CFILT Social Media outlets (LinkedIn, Twitter and Instagram). Increased online visibility by **619%** on LinkedIn over the duration.

Lead Teaching Assistant | Advanced-NLP | Qualcomm Apr'23 – Sept'23

- Helped in delivering a course on NLP to Qualcomm Engineers. Tasks handled: Preparing the Curriculum, Coding Demos and Preparing Quizzes.

TECHNICAL SKILLS

Languages C, Java, Python, C++, Bash, Javascript

Tools & Frameworks PyTorch, JAX, Flax, TensorFlow, HuggingFace, Ray, DeepSpeed, Streamlit, Gradio, Spring